



# Sentiment Analysis for Marketing


## PHASE -4


DATE	24 October 2023
TEAM ID	Proj-212173-Team-1
PROJECT NAME	Sentiment analysis for Marketing

```
✓ [13] MODEL = 'bert-base-uncased'
1s tokenizer = BertTokenizer.from_pretrained(MODEL, do_lower_case=True)
```

Downloading (...)okenizer\_config.json: 100%  28.0/28.0 [00:00<00:00, 664B/s]


Downloading (...)solve/main/vocab.txt: 100%  232k/232k [00:00<00:00, 819kB/s]

Downloading (...)main/tokenizer.json: 100%  466k/466k [00:00<00:00, 12.5MB/s]

Downloading (...)lve/main/config.json: 100%  570/570 [00:00<00:00, 19.6kB/s]

BERT base model (uncased) Pretrained model on English language using a masked language modeling (MLM) objective. It was introduced in this paper and first released in this repository. This model is uncased: it does not make a difference between english and English.

```
✓ [14] sns.displot(tokens)
8s plt.xlim([0, 60]);
plt.xlabel('Token count');
```



Token Count means the number of Tokens required to operate a single instance of a Software product in accordance with the provisions of this Agreement.

```

✓ 0s from torch.utils.data import Dataset, DataLoader

# Define a custom dataset, more info on how to build custom dataset can be
# found at https://pytorch.org/tutorials/beginner/data\_loading\_tutorial.html
class CustomDataset(Dataset):

    def __init__(
        self,
        tweets,
        labels,
        tokenizer,
        max_length
    ):
        self.tweets = tweets
        self.labels = labels
        self.tokenizer = tokenizer
        self.max_length = max_length

    def __len__(self):
        return len(self.tweets)

    def __getitem__(self, idx):
        tweet = self.tweets[idx]
        label = self.labels[idx]

        tokenize = self.tokenizer.encode_plus(
            tweet,
            add_special_tokens=True,
            max_length=self.max_length,
            return_token_type_ids=False,
            padding='max_length',
            return_attention_mask=True,
            return_tensors='pt'
        )
        return {
            'tweet': tweet,
            'input_ids': tokenize['input_ids'].flatten(),
            'attention_mask': tokenize['attention_mask'].flatten(),
            'targets': torch.tensor(label, dtype=torch.long)}

```

In the above program return the values of tweet,input\_ids,attention\_mask and target to the program.It is the main part of import the dataset and create the class then define the functions and finally produce the output.

```
✓ 0s ▶ MAX_LENGTH = 64
TEST_SIZE = 0.1
VALID_SIZE = 0.5
BATCH_SIZE = 16
NUM_WORKERS = 2

train_sampler, test_sampler = train_test_split(df, test_size=TEST_SIZE, random_state=RANDOM_STATE)
valid_sampler, test_sampler = train_test_split(test_sampler, test_size=VALID_SIZE, random_state=RANDOM_STATE)

train_set = CustomDataset(
    train_sampler['text'].to_numpy(),
    train_sampler['labels'].to_numpy(),
    tokenizer,
    MAX_LENGTH
)
test_set = CustomDataset(
    test_sampler['text'].to_numpy(),
    test_sampler['labels'].to_numpy(),
    tokenizer,
    MAX_LENGTH
)
valid_set = CustomDataset(
    valid_sampler['text'].to_numpy(),
    valid_sampler['labels'].to_numpy(),
    tokenizer,
    MAX_LENGTH
)
```

```
train_loader = DataLoader(train_set, batch_size=BATCH_SIZE, num_workers=NUM_WORKERS)
test_loader = DataLoader(test_set, batch_size=BATCH_SIZE, num_workers=NUM_WORKERS)
valid_loader = DataLoader(valid_set, batch_size=BATCH_SIZE, num_workers=NUM_WORKERS)
```

The motivation is quite simple: you should separate your data into train, validation, and test splits to prevent your model from overfitting and to accurately evaluate your model. Accurate training data helps the model learn the right patterns, validation data helps developers fine-tune the model correctly, and test data provides trustworthy metrics so they can confidently deploy their AI solution.

✓  
0s



```
from torch import nn
class AirlineSentimentClassifier(nn.Module):

    def __init__(self, num_labels):
        super(AirlineSentimentClassifier, self).__init__()
        self.bert = BertModel.from_pretrained(MODEL)
        self.dropout = nn.Dropout(p=0.2)
        self.classifier = nn.Linear(self.bert.config.hidden_size, num_labels)

    def forward(self, input_ids, attention_mask):
        outputs = self.bert(
            input_ids=input_ids,
            attention_mask=attention_mask
        )
        pooled_output = outputs[1]
        pooled_output = self.dropout(pooled_output)
        out = self.classifier(pooled_output)
        return out
```

✓  
6s

```
[19] model = AirlineSentimentClassifier(len(labels_map))
print(model)

# Move tensors to GPU on CUDA enables devices
if device:
    model.cuda()
```

Downloading model.safetensors: 100%  440M/440M [00:04<00:00, 33.2MB/s]

```
AirlineSentimentClassifier(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
    )
  )
```

```

        (intermediate): BertIntermediate(
          (dense): Linear(in_features=768, out_features=3072, bias=True)
          (intermediate_act_fn): GELUActivation()
        )
        (output): BertOutput(
          (dense): Linear(in_features=3072, out_features=768, bias=True)
          (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
      )
    )
    (pooler): BertPooler(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (activation): Tanh()
    )
  )
  (dropout): Dropout(p=0.2, inplace=False)
  (classifier): Linear(in_features=768, out_features=3, bias=True)
)

```

CUDA® is a parallel computing platform and programming model invented by NVIDIA. It enables dramatic increases in computing performance by harnessing the power of the graphics processing unit (GPU).



```

15 n_epochs = 10
   learning_rate = 2e-5

   # Loss function
   criterion = nn.CrossEntropyLoss()

   # Optimizer
   optimizer = AdamW(model.parameters(), lr=learning_rate, correct_bias=False)

   # Define scheduler
   training_steps = len(train_loader)*n_epochs
   scheduler = get_linear_schedule_with_warmup(
       optimizer,
       num_warmup_steps=0,
       num_training_steps=training_steps
   )

   /usr/local/lib/python3.10/dist-packages/transformers/optimization.py:411: FutureWarning: This implementation of AdamW is de
warnings.warn(

```

In machine learning, a loss function and an optimizer are two essential components that help to improve the performance of a model. A loss function measures the difference between the predicted output of a model and the actual output, while an optimizer adjusts the model's parameters to minimize the loss function.

```

# Track changes in validation loss
valid_loss_min = np.Inf

for epoch in range(1, n_epochs+1):

    # Setting training and validation loss
    train_loss = []
    validation_loss = []
    tr_predictions = 0
    acc = 0
    val_predictions = 0

    #####
    # Train the model #
    #####
    model = model.train()
    for data in train_loader:

        # Moving tensors to GPU on CUDA enabled devices
        if device:
            input_ids, attention_mask, targets = data["input_ids"].cuda(), data["attention_mask"].cuda(), data["targets"]
        # Clear the gradients of variables
        optimizer.zero_grad()

        #### Forward pass
        # Pass input through the model
        output = model(
            input_ids=input_ids,
            attention_mask=attention_mask
        )
        # Compute batch loss
        loss = criterion(output, targets)
        # Convert output probabilities to class probabilities
        _, pred = torch.max(output, 1)
        # Track correct predictions
        tr_predictions += torch.sum(pred == targets)

        #### Backward Pass
        # Compute gradients wrt to model parameters
        loss.backward()
        # To avoid exploding gradients, we clip the gradients of the model
        nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
        # Perform parameter update
        optimizer.step()
        # Update learning rate
        scheduler.step()
        # Update loss per mini batches
        train_loss.append(loss.item())

    #####
    # Validate the model #
    #####
    model.eval()
    with torch.no_grad():
        for data in valid_loader:

```



```

# Moving tensors to GPU on CUDA enabled devices
if device:
    input_ids, attention_mask, targets = data["input_ids"].cuda(), data["attention_mask"].cuda(), data["tar

#### Forward pass
# Pass input through the model
output = model(
    input_ids=input_ids,
    attention_mask=attention_mask
)
# Compute batch loss
loss = criterion(output, targets)
# Convert output probabilities to class probabilities
_, pred = torch.max(output, 1)
# Update loss per mini batches
validation_loss.append(loss.item())
# Track correct predictions
val_predictions += torch.sum(pred == targets)

```

```

# Compute accuracy
train_accuracy = tr_predictions.double()/len(train_sampler)
val_accuracy = val_predictions.double()/len(valid_sampler)

# Print loss statistics
print('Epoch: {}/{} \n\tTraining Loss: {:.6f} \n\tValidation Loss: {:.6f} \n\tTrain Accuracy: {:.6f} \n\tVal Accura

# Save model if validation loss is decreased
if val_accuracy > acc:
    print('Saving model...')
    torch.save(model.state_dict(), 'bert_base_fine_tuned.pt')
    acc = val_accuracy

```

One of the most widely used metrics combinations is training loss + validation loss over time. The training loss indicates how well the model is fitting the training data, while the validation loss indicates how well the model fits new data.

At times, the validation loss is greater than the training loss. This may indicate that the model is underfitting. Underfitting occurs when the model is unable to accurately model the training data, and hence generates large errors.

Typically validation loss should be similar to but slightly higher than training loss. As long as validation loss is lower than or even equal to training loss one should keep doing more training.



```
Epoch: 1/10
  Training Loss: 0.478485
  Validation Loss: 0.426510
  Train Accuracy: 0.813221
  Val Accuracy: 0.848361
Saving model...
Epoch: 2/10
  Training Loss: 0.251598
  Validation Loss: 0.587404
  Train Accuracy: 0.912720
  Val Accuracy: 0.837432
Saving model...
Epoch: 3/10
  Training Loss: 0.147462
  Validation Loss: 0.694001
  Train Accuracy: 0.958333
  Val Accuracy: 0.848361
Saving model...
Epoch: 4/10
  Training Loss: 0.095958
  Validation Loss: 0.852052
  Train Accuracy: 0.976548
  Val Accuracy: 0.841530
Saving model...
Epoch: 5/10
  Training Loss: 0.062927
  Validation Loss: 0.967488
  Train Accuracy: 0.985504
  Val Accuracy: 0.842896
Saving model...
Epoch: 6/10
  Training Loss: 0.042360
  Validation Loss: 1.066000
  Train Accuracy: 0.990437
  Val Accuracy: 0.840164
Saving model...
Epoch: 7/10
  Training Loss: 0.032142
  Validation Loss: 1.132496
  Train Accuracy: 0.992410
  Val Accuracy: 0.833333
Saving model...
Epoch: 8/10
  Training Loss: 0.024429
  Validation Loss: 1.184951
  Train Accuracy: 0.993777
  Val Accuracy: 0.829235
Saving model...
Epoch: 9/10
  Training Loss: 0.018996
  Validation Loss: 1.230268
  Train Accuracy: 0.994991
  Val Accuracy: 0.831967
Saving model...
Epoch: 10/10
  Training Loss: 0.015075
  Validation Loss: 1.244014
  Train Accuracy: 0.995826
  Val Accuracy: 0.830601
Saving model...
```

when you see that the model performs well on the training data but does not perform well on the evaluation data. This is because the model is memorizing the data it has seen and is unable to generalize to unseen examples.

```

# Track test loss
test_loss = 0.0
class_predictions = list(0. for i in range(3))
class_total = list(0. for i in range(3))
predictions = []
labels = []

model.eval()
with torch.no_grad():
    for data in test_loader:

        # Moving tensors to GPU on CUDA enabled devices
        if device:
            input_ids, attention_mask, targets = data["input_ids"].cuda(), data["attention_mask"].cuda(), data["targets"]

        #### Forward pass
        # Pass input through the model
        output = model(
            input_ids=input_ids,
            attention_mask=attention_mask
        )
        # Compute batch loss
        loss = criterion(output, targets)
        # Update loss
        test_loss += loss.item()
        # convert output probabilities to predicted class
        _, pred = torch.max(output, 1)

        predictions.extend(pred)
        labels.extend(targets)

predictions = torch.stack(predictions) if not device else torch.stack(predictions).cpu()
labels = torch.stack(labels) if not device else torch.stack(labels).cpu()

```

```

print(classification_report(predictions, labels, target_names=['neutral', 'positive', 'negative']))

```


	precision	recall	f1-score	support
neutral	0.68	0.71	0.69	146
positive	0.79	0.81	0.80	118
negative	0.92	0.90	0.91	468
accuracy			0.85	732
macro avg	0.79	0.81	0.80	732
weighted avg	0.85	0.85	0.85	732

```

cm = confusion_matrix(labels, predictions)
heatmap = sns.heatmap(cm, annot=True, fmt='d', cmap='Greens')
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=0, ha='right')
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=30, ha='right')
plt.xlabel('True sentiment')
plt.ylabel('Predicted sentiment');

```

The phrase 'the true sentiment' is correct and usable in written English. You can use it to refer to an honest or heartfelt emotion that someone is feeling. For example, "Let us honor the true sentiment behind the movement."



1. Knowledge Based: This approach includes the classification of text based on words that are associated with emotion.
2. Statistical: This approach utilizes machine learning algorithms such as latent semantic analysis and deep learning for accurate detection of sentiment.